

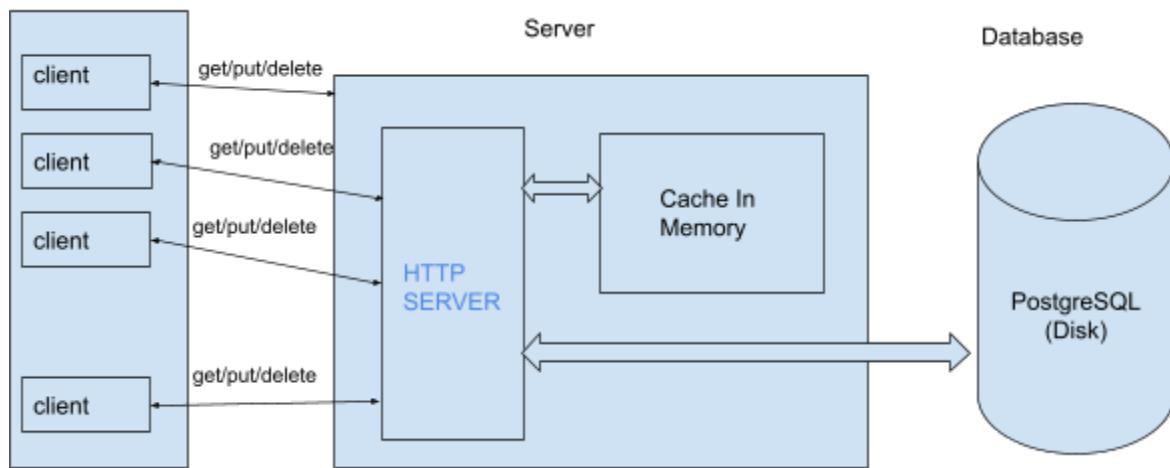
# HTTP-based Key-Value Server

Roll no- 25m0750

Paper code-CS744 Autumn 2025

## Description:

A high-performance, multi-threaded **HTTP Key-Value Store** in C++. It uses **PostgreSQL** for persistent storage and an in-memory **LRU (Least Recently Used) Cache** to speed up data retrieval.



## 1. Architectural Overview

The server acts as a middleware between HTTP clients and a PostgreSQL database.

- **Clients** send HTTP requests (GET, PUT, DELETE) to endpoints like `http://localhost:8080/kv/myKey`.
- **Server** checks its fast in-memory cache first.
- **Database** (PostgreSQL) is used for persistent data.

## 2. Core Components

### A. PGPool (Database Connection Pool)

Connecting to a database is expensive and slow. This class maintains a pool of open connections to PostgreSQL to be reused by different threads.

- **Initialization:** On startup, it opens pool\_size (default 4) connections.
- **Preparation:** It creates the kv\_store table if it doesn't exist and "prepares" SQL statements (kv\_get, kv\_put, kv\_del) for efficiency.
- **Thread Safety:** Uses a mutex and condition\_variable to allow multiple HTTP worker threads to safely acquire() and release() connections.

## B. LRUCache (In-Memory Cache)

This is a classic implementation of a Least Recently Used cache. It keeps the most frequently accessed data in RAM for instant retrieval.

- **Mechanism:** It uses two data structures simultaneously:
  1. std::list: Maintains the order of use. The front is the Most Recently Used (MRU), and the back is the Least Recently Used (LRU).
  2. std::unordered\_map: Provides O(1) fast lookups, mapping a "key" to an iterator pointing to that item's location in the std::list.
- **Eviction Policy:** When the cache is full and a new item is added, it removes the item at the back of the list (the LRU item) to make space.
- **Thread Safety:** Protected by a std::mutex because multiple HTTP threads might try to read/update the cache at the same time.

## C. KVHandler (HTTP Request Handler)

This class connects the web server (CivetServer), the database (PGPool), and the cache (LRUCache).

- **Cache Warming:** When the server starts, warmUpCache() queries the *entire* database and loads it into the cache (up to the cache's max size). This ensures the cache is useful immediately upon startup.
  - **Request Routing:** It parses URI paths (e.g., extracting myKey from /kv/myKey) and routes them to doGet, doPut, or doDelete.
- 

## 3. How Caching is Implemented

The code uses standard caching strategies to ensure data consistency between memory and the database.

### Read Path (GET Request) - "Cache-Aside" / "Read-Through"

When a client requests a key:

1. **Check Cache:** The handler first asks the LRUCache.
  - o *HIT*: If found, returns the value immediately (super fast). The cache marks this item as "most recently used."
2. **Check DB (Miss):** If not in cache, it acquires a DB connection and queries PostgreSQL.
  - o *Not found in DB*: Returns 404.
  - o *Found in DB*:
    1. Value is fetched from DB.
    2. Value is **inserted into the Cache** so future requests for this key will HIT.
    3. Value is returned to the client.

### **Write Path (PUT Request) - "Write-Through"**

When a client updates or inserts a key:

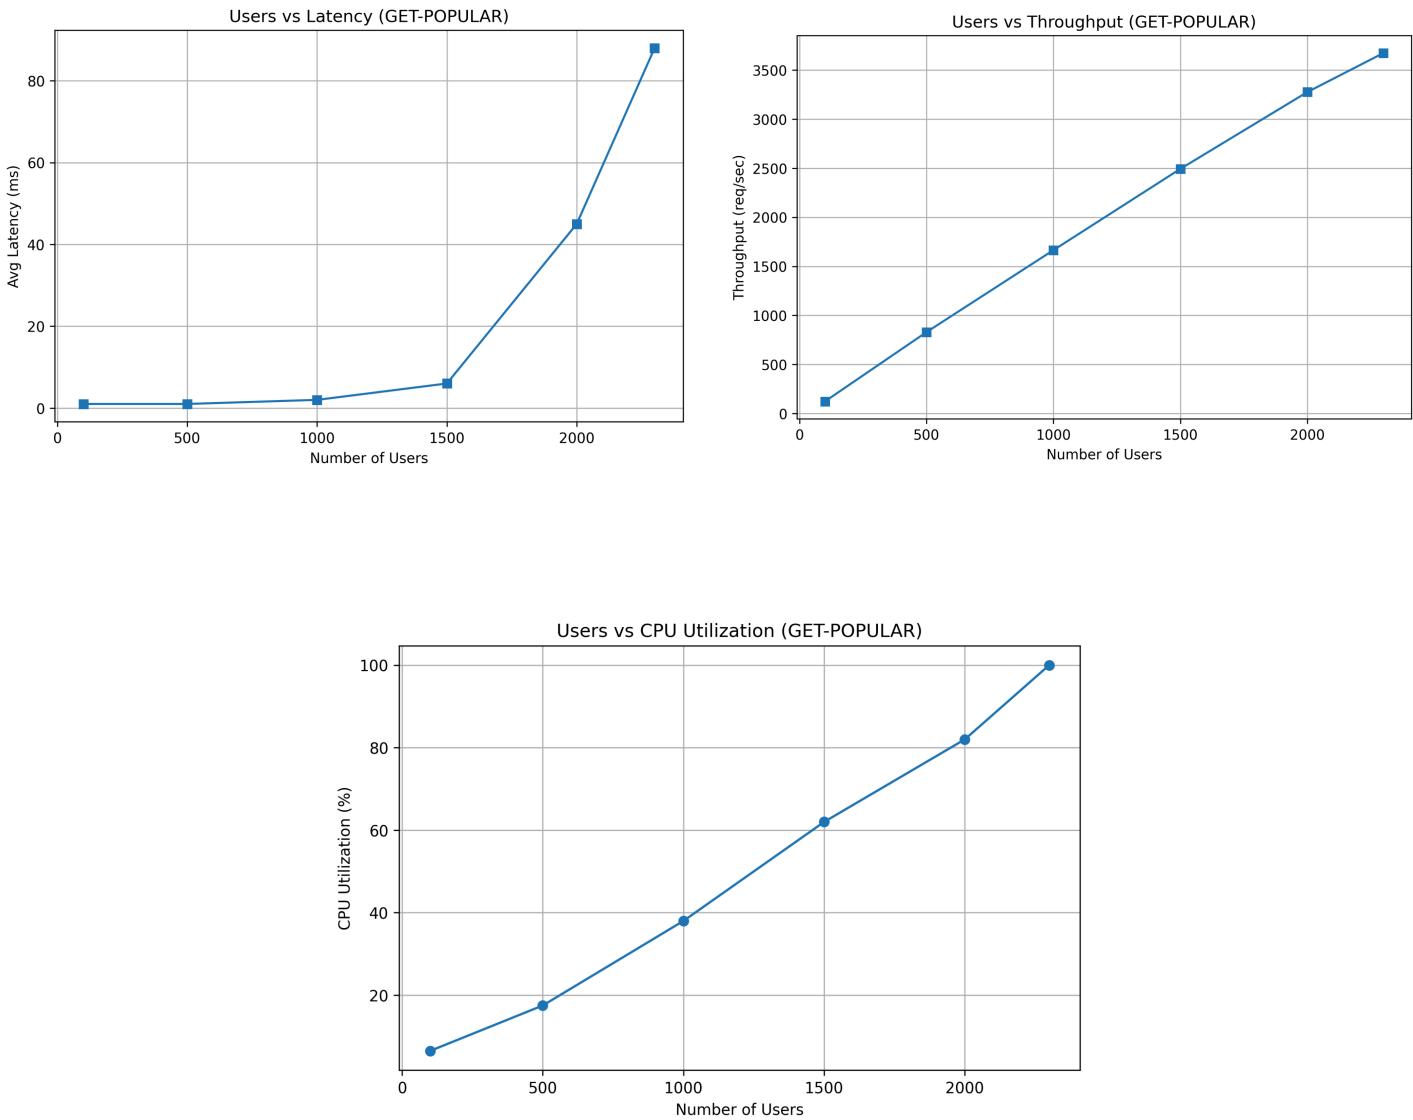
1. **Write to DB:** It first persists the data to PostgreSQL to ensure durability.
2. **Update Cache:** If the DB write succeeds, it immediately updates (put) the item in the LRUCache. This ensures subsequent reads get the new value.

### **Delete Path (DELETE Request)**

1. **Delete from DB:** Removes the row from PostgreSQL.
2. **Invalidate Cache:** Calls `cache_.erase(key)` to remove it from memory, ensuring no one reads stale data.

Results:

## Workload GET-popular:

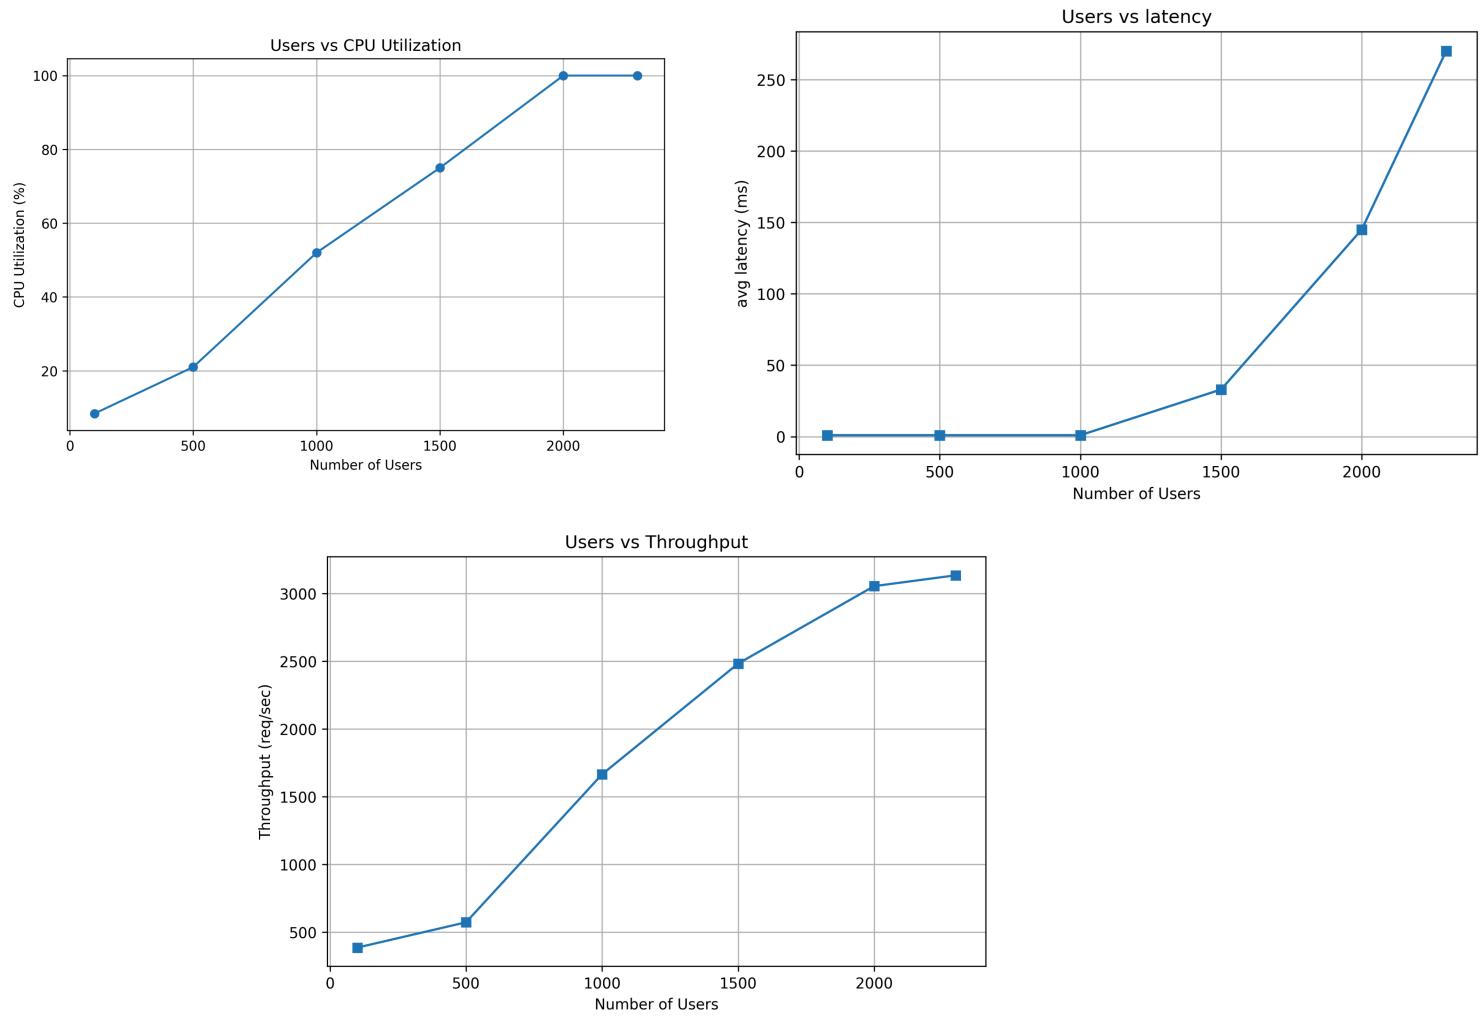


This workload access data from cache only , no io over head.  
For this workload cpu in server is bottleneck

### CPU Utilization Trend

- CPU increases steadily from 8% → 100%.
- System becomes fully saturated at 2300 users

Workload get-all:

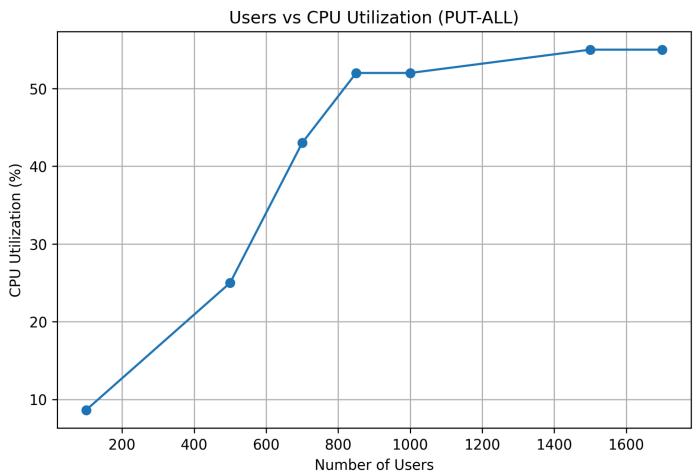
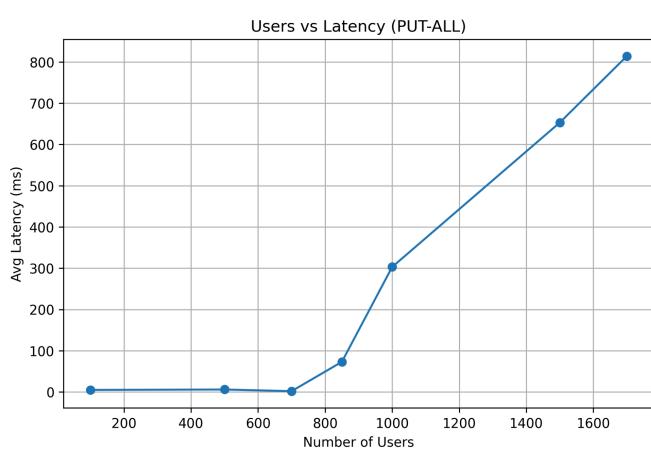
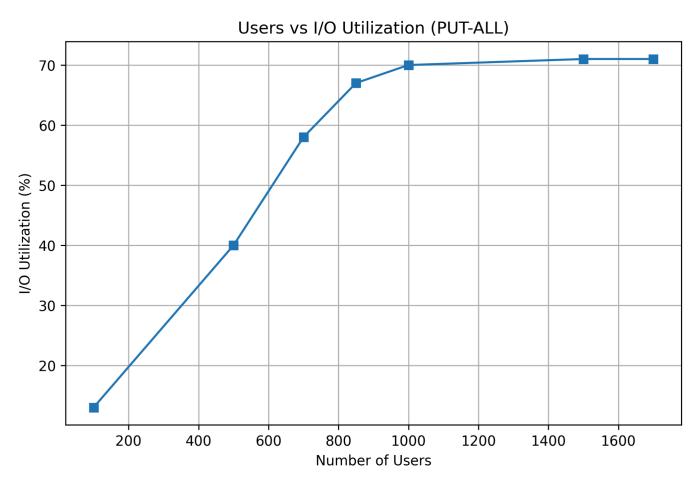
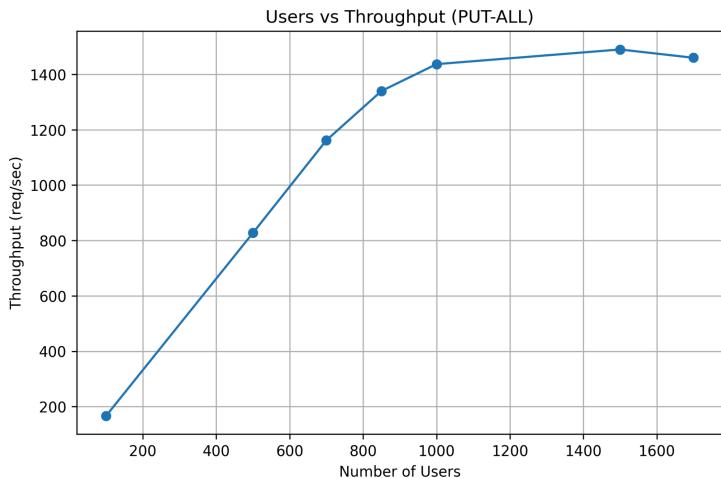


This workload simulates get request from cache if miss then request served by database , only read overhead. But no write in db table;

### CPU Utilization Trend

- CPU increases steadily from 8% → 100%.
- System becomes fully saturated at 2000 users

## Workload put\_all:



This workload simulates only put request ,served by database , read overhead for finding the key and write value on db table;

### CPU Utilization Trend

- CPU is not the bottleneck as only 50-55% is utilized.
- I/O operation in database is main bottleneck here. After 70 % utilization saturates its capacity.

Github Link:

[https://github.com/Blooming2003sun/decs\\_kvserver](https://github.com/Blooming2003sun/decs_kvserver)